

# Implement a CNN Using Keras

# 1. Load and Process the MNIST Dataset

## Load data

```
from keras.datasets import mnist

(x_train, y_train), (x_test, y_test) = mnist.load_data()

print('Shape of x_train: ' + str(x_train.shape))
print('Shape of x_test: ' + str(x_test.shape))
print('Shape of y_train: ' + str(y_train.shape))
print('Shape of y_test: ' + str(y_test.shape))
```

Shape of x\_train: (60000, 28, 28)

Shape of x\_test: (10000, 28, 28)

Shape of y\_train: (60000,)

Shape of y\_test: (10000,)

# 1. Load and Process the MNIST Dataset

**Reshape:** convert the  $60000 \times 28 \times 28$  dataset to a  $60000 \times 28 \times 28 \times 1$  tensor.

```
x_train_vec = x_train.reshape((60000, 28, 28, 1)) / 255.0
x_test_vec = x_test.reshape((10000, 28, 28, 1)) / 255.0

print('Shape of x_train_vec is ' + str(x_train_vec.shape))

Shape of x_train_vec is (60000, 28, 28, 1)
```

# 1. Load and Process the MNIST Dataset

**One-hot encode:** convert the **labels** (scalars in  $\{0, 1, \dots, 9\}$ ) to 10-dim vectors.

```
def to_one_hot(labels, dimension=10):
    results = np.zeros((len(labels), dimension))
    for i, label in enumerate(labels):
        results[i, label] = 1.
    return results

y_train_vec = to_one_hot(y_train)
y_test_vec = to_one_hot(y_test)
print('Shape of y_train_vec is ' + str(y_train_vec.shape))
```

Shape of y\_train\_vec is (60000, 10)

# 1. Load and Process the MNIST Dataset

Partition the **training** set to **training** and **validation** sets.

```
# Partition to training and validation sets

rand_indices = np.random.permutation(60000)
train_indices = rand_indices[0:50000]
valid_indices = rand_indices[50000:60000]

x_valid_vec = x_train_vec[valid_indices, :, :, :]
y_valid_vec = y_train_vec[valid_indices, :]

x_train_vec = x_train_vec[train_indices, :, :, :]
y_train_vec = y_train_vec[train_indices, :]

print('Shape of x_valid_vec: ' + str(x_valid_vec.shape))
print('Shape of y_valid_vec: ' + str(y_valid_vec.shape))
print('Shape of x_train_vec: ' + str(x_train_vec.shape))
print('Shape of y_train_vec: ' + str(y_train_vec.shape))
```

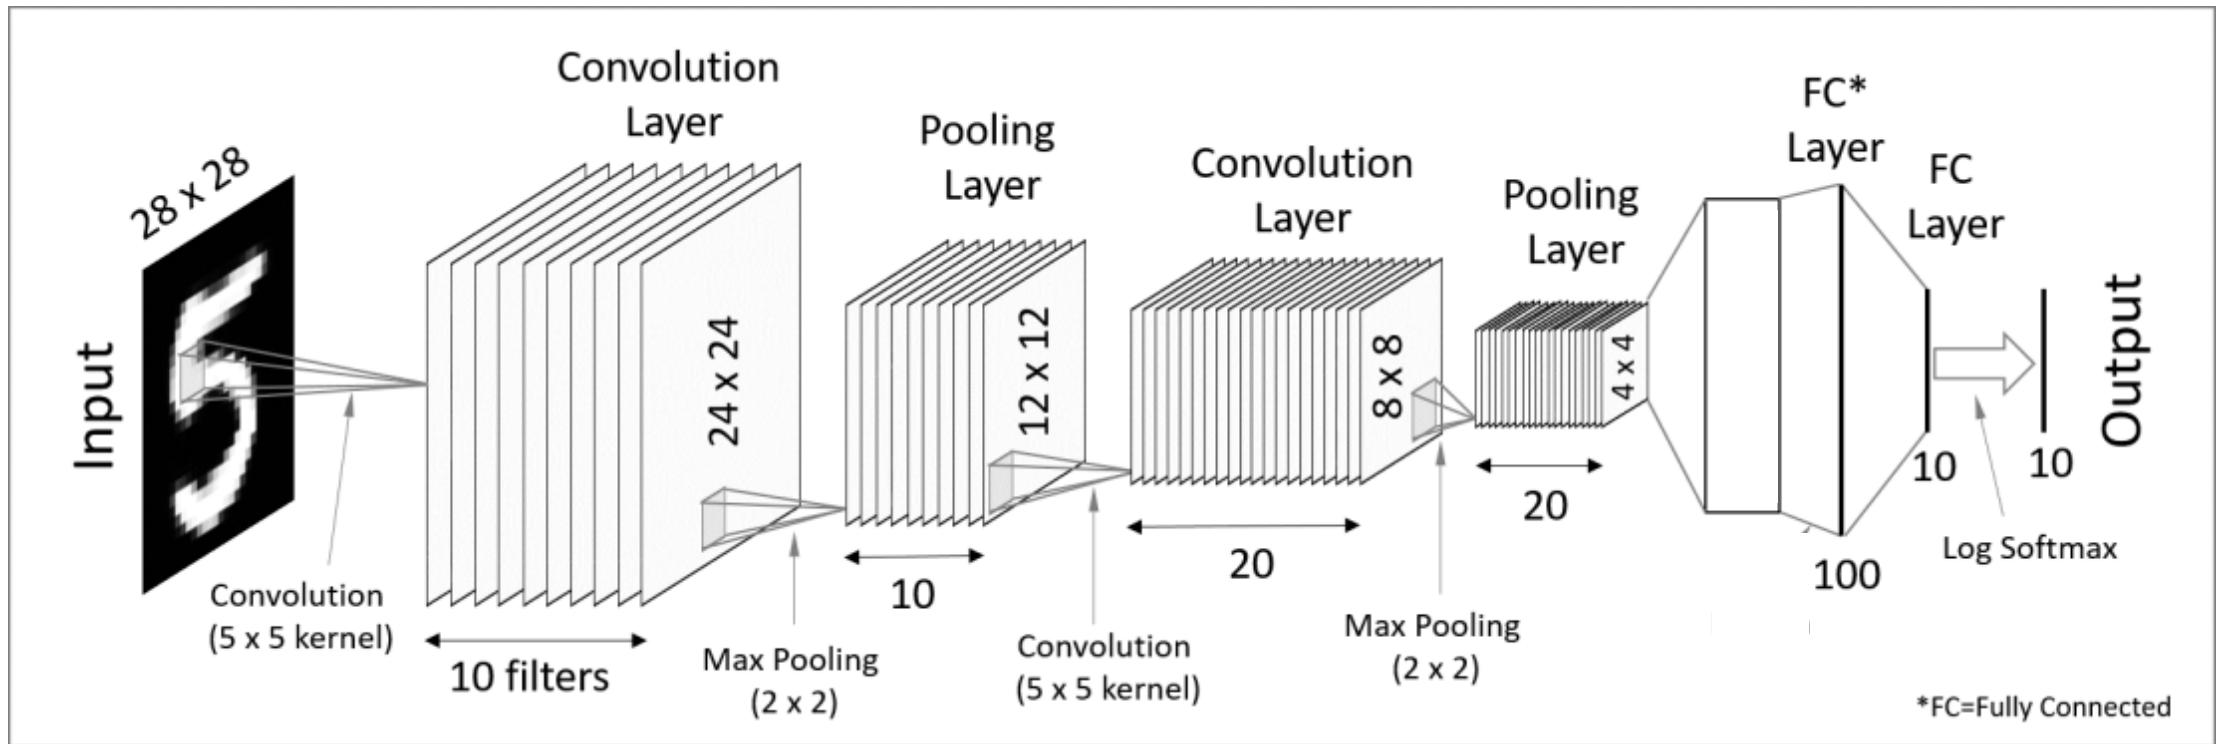
Shape of x\_valid\_vec: (10000, 28, 28, 1)

Shape of y\_valid\_vec: (10000, 10)

Shape of x\_train\_vec: (50000, 28, 28, 1)

Shape of y\_train\_vec: (50000, 10)

## 2. Build the CNN



## 2. Build the CNN

```
from keras import models
from keras import layers

model = models.Sequential()
model.add(layers.Conv2D(10, (5, 5), activation='relu', input_shape=(28, 28, 1)))
```

Filter number      Filter shape



- Default
- Stride: 1
  - Zero-padding: no padding

## 2. Build the CNN

```
from keras import models
from keras import layers

model = models.Sequential()
model.add(layers.Conv2D(10, (5, 5), activation='relu', input_shape=(28, 28, 1)))
model.add(layers.MaxPooling2D((2, 2)))
```

pool\_size



Default

- Stride: pool\_size

## 2. Build the CNN

```
from keras import models
from keras import layers

model = models.Sequential()
model.add(layers.Conv2D(10, (5, 5), activation='relu', input_shape=(28, 28, 1)))
model.add(layers.MaxPooling2D((2, 2))) Output: 12×12×10
model.add(layers.Conv2D(20, (5, 5), activation='relu')) Output: 8×8×20
model.add(layers.MaxPooling2D((2, 2))) Output: 4×4×20
```

## 2. Build the CNN

```
from keras import models
from keras import layers

model = models.Sequential()
model.add(layers.Conv2D(10, (5, 5), activation='relu', input_shape=(28, 28, 1)))
model.add(layers.MaxPooling2D((2, 2))) Output: 12×12×10
model.add(layers.Conv2D(20, (5, 5), activation='relu')) Output: 8×8×20
model.add(layers.MaxPooling2D((2, 2))) Output: 4×4×20
model.add(layers.Flatten()) Output: 320
```

## 2. Build the CNN

```
from keras import models
from keras import layers

model = models.Sequential()
model.add(layers.Conv2D(10, (5, 5), activation='relu', input_shape=(28, 28, 1)))
model.add(layers.MaxPooling2D((2, 2))) Output: 12×12×10
model.add(layers.Conv2D(20, (5, 5), activation='relu')) Output: 8×8×20
model.add(layers.MaxPooling2D((2, 2))) Output: 4×4×20
model.add(layers.Flatten()) Output: 320
model.add(layers.Dense(100, activation='relu')) Output: 100
model.add(layers.Dense(10, activation='softmax')) Output: 10
```

## 2. Build the CNN

```
# print the summary of the model.  
model.summary()
```

Layer (type)	Output Shape	Param #
=====		
conv2d_1 (Conv2D)	(None, 24, 24, 10)	260
max_pooling2d_1 (MaxPooling2D)	(None, 12, 12, 10)	0
conv2d_2 (Conv2D)	(None, 8, 8, 20)	5020
max_pooling2d_2 (MaxPooling2D)	(None, 4, 4, 20)	0
flatten_1 (Flatten)	(None, 320)	0
dense_1 (Dense)	(None, 100)	32100
dense_2 (Dense)	(None, 10)	1010
=====		
Total params: 38,390		
Trainable params: 38,390		
Non-trainable params: 0		

### 3. Train the CNN

Specify: optimization algorithm, learning rate (LR), loss function, and metric.

```
from keras import optimizers  
model.compile(optimizers.RMSprop(lr=0.0001),  
              loss='categorical_crossentropy',  
              metrics=['accuracy'])
```

### 3. Train the CNN

**Specify:** batch size and number of epochs.

```
history = model.fit(x_train_vec, y_train_vec,  
                     batch_size=128, epochs=50,  
                     validation_data=(x_valid_vec, y_valid_vec))
```

```
Train on 50000 samples, validate on 10000 samples  
Epoch 1/50  
50000/50000 [=====] - 12s 239us/step - loss: 1.2455 - acc: 0.7201 - val_loss: 0.4805 - val_acc: 0.8707  
Epoch 2/50  
50000/50000 [=====] - 12s 235us/step - loss: 0.3563 - acc: 0.8995 - val_loss: 0.2851 - val_acc: 0.9172  
Epoch 3/50  
50000/50000 [=====] - 12s 244us/step - loss: 0.2474 - acc: 0.9271 - val_loss: 0.2230 - val_acc: 0.9365
```

●  
●  
●

```
Epoch 49/50  
50000/50000 [=====] - 12s 240us/step - loss: 0.0244 - acc: 0.9924 - val_loss: 0.0439 - val_acc: 0.9875  
Epoch 50/50  
50000/50000 [=====] - 12s 239us/step - loss: 0.0238 - acc: 0.9927 - val_loss: 0.0446 - val_acc: 0.9881
```

# 4. Examine the Results

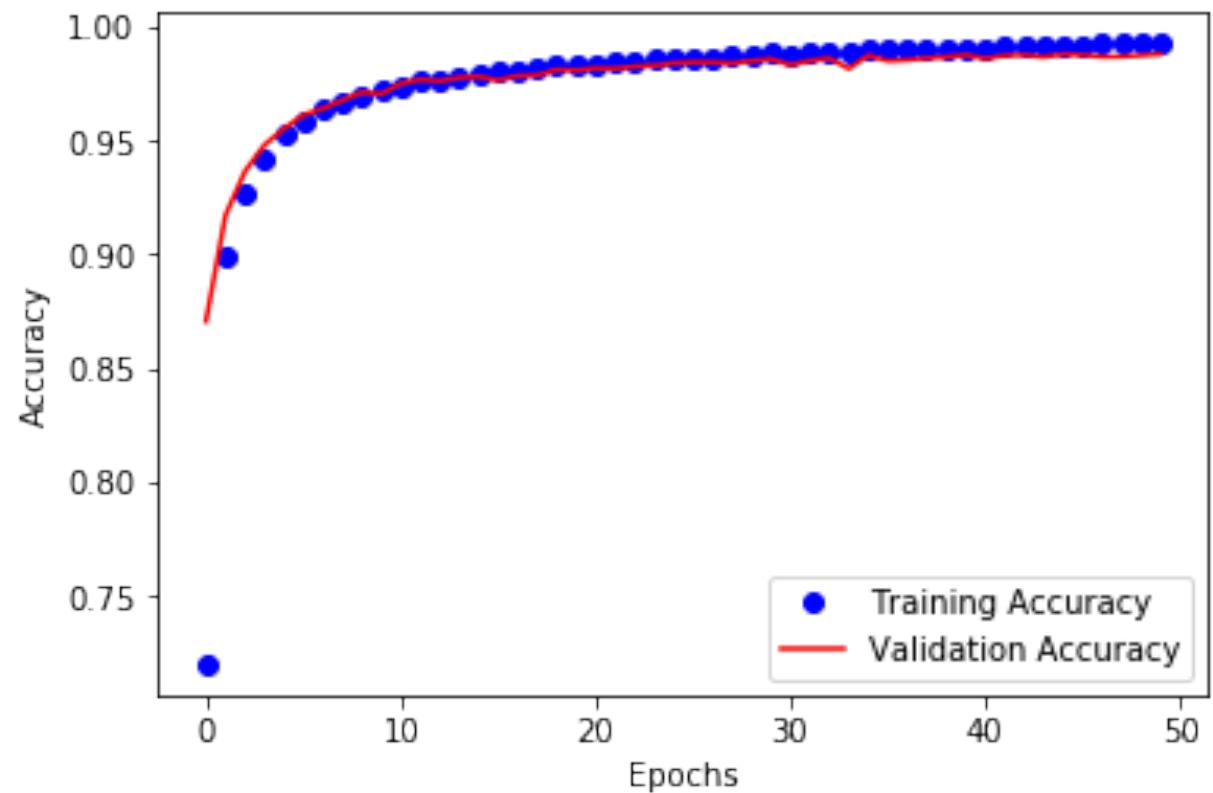
Plot the *accuracy* against *epochs* (1 epoch = 1 pass over the data).

```
import matplotlib.pyplot as plt
%matplotlib inline

epochs = range(50) # 50 is the number of epochs
train_acc = history.history['acc']
valid_acc = history.history['val_acc']
plt.plot(epochs, train_acc, 'bo', label='Training Accuracy')
plt.plot(epochs, valid_acc, 'r', label='Validation Accuracy')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.legend()
plt.show()
```

# 4. Examine the Results

Plot the *accuracy* against *epochs* (1 epoch = 1 pass over the data).



Why using the validation set?

- Tune the hyper-parameters, e.g., learning rate, filter shape, filter number, etc.
- Check overfitting or underfitting.

## 4. Examine the Results

Evaluate the model on the ***test*** set. (So far, the model has not seen the ***test*** set.)

```
loss_and_acc = model.evaluate(x_test_vec, y_test_vec)
print('loss = ' + str(loss_and_acc[0]))
print('accuracy = ' + str(loss_and_acc[1]))
```

```
10000/10000 [=====] - 1s 127us/step
loss = 0.03570800104729715
accuracy = 0.9883
```

## 4. Examine the Results

- Training accuracy: 99.3%
- Validation accuracy: 98.8%
- Test accuracy: 98.8%

# Summary

# 1. Convolution

- Matrix convolutional and tensor convolution.
- Concepts:
  - filter (kernel)
  - patch
  - output (feature map)
  - zero-padding
  - stride

## 2. Convolutional Neural Network

- Convolutional layers.
  - Filter number.
  - Filter shape.
  - Stride (default: one).
  - Zero-padding (default: no padding).

## 2. Convolutional Neural Network

- Convolutional layers.
  - Filter number.
  - Filter shape.
  - Stride (default: one).
  - Zero-padding (default: no padding).
- Activation functions.
- Pooling layers.
  - MaxPool, AveragePool, etc.
  - Pool size.
  - Pool stride (default: pool size).

### 3. Build and Train CNN using Keras

- Build a network.
  - Add Conv, Pool, FC layers.
- Compile the model.
  - Specify optimization algorithm, loss function, and metrics.
- Fit the model on training data.

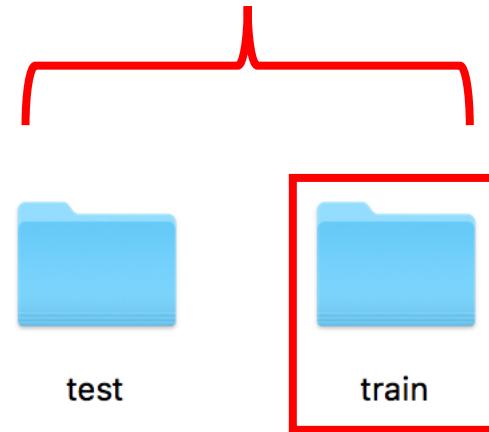
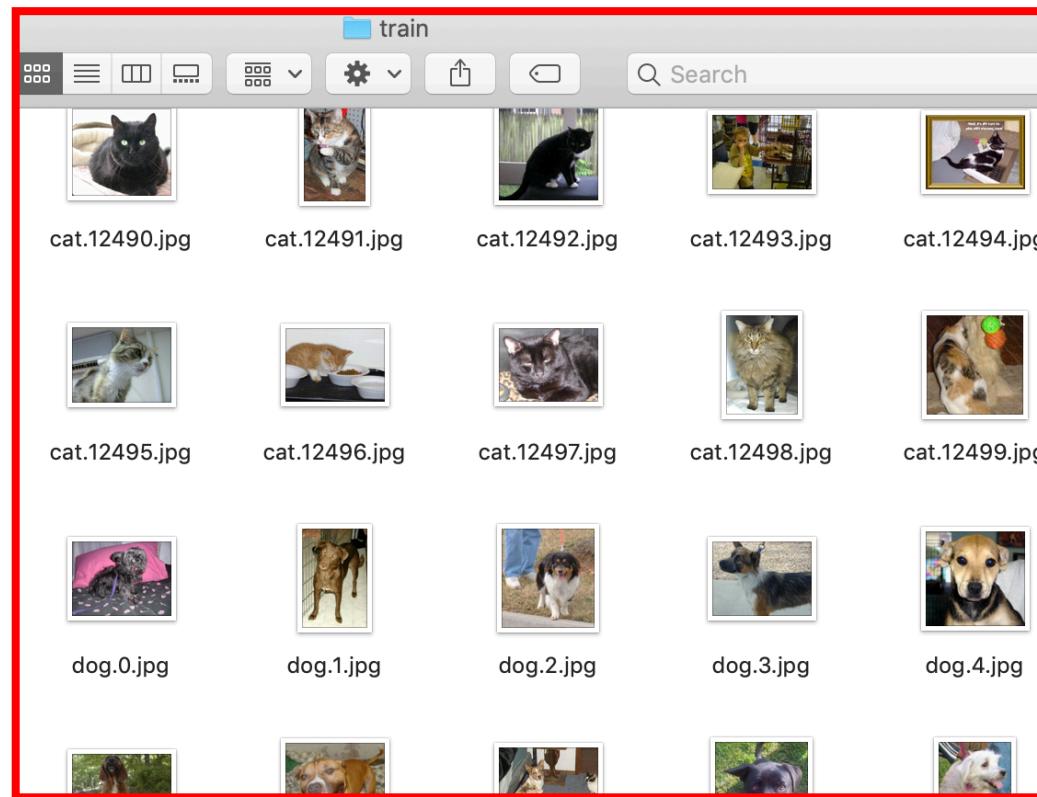
# The Dogs vs. Cats Dataset

# The Dogs vs. Cats Dataset



# The Dogs vs. Cats Dataset

Download: <https://www.kaggle.com/c/dogs-vs-cats/data>



# The Dogs vs. Cats Dataset

Download: <https://www.kaggle.com/c/dogs-vs-cats/data>



cat.0.jpg



cat.1.jpg



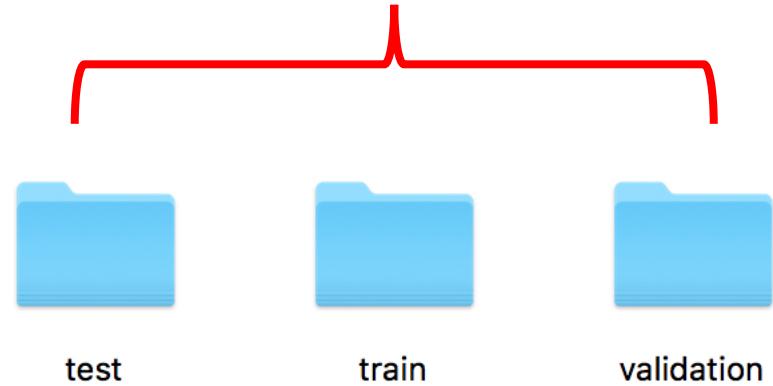
cat.2.jpg



cat.3.jpg



cat.4.jpg



cat.5.jpg



cat.6.jpg



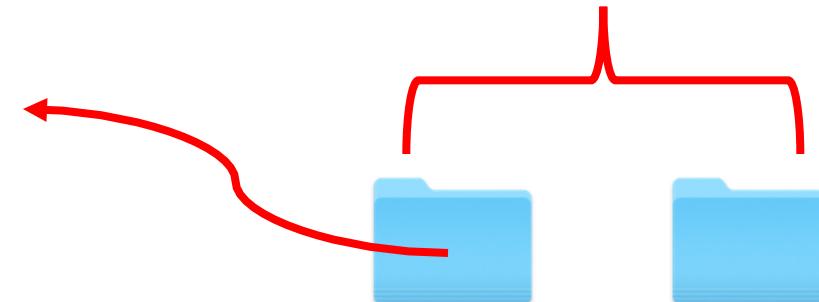
cat.7.jpg



cat.8.jpg



cat.9.jpg



cat.10.jpg



cat.11.jpg



cat.12.jpg



cat.13.jpg



cat.14.jpg

# The Dogs vs. Cats Dataset

- The dataset has 25,000 training samples.
- I use a subset:
  - 2000 images for training,
  - 1000 images for validation,
  - 1000 images for test.

# Implement a CNN Using Keras

# 1. Load and Process the Dataset

- Currently, the data sit on a drive as JPEG files.
- Data processing:
  1. Read the picture files.
  2. Decode the JPEG content to order-3 tensors.
  3. Resize the images to the the same shape, e.g.,  $150 \times 150 \times 3$ .
  4. Rescale the pixel values (between 0 and 255) to the  $[0, 1]$  interval.

# 1. Load and Process the Dataset

```
from keras.preprocessing.image import ImageDataGenerator  
  
# All images will be rescaled by 1./255  
train_datagen = ImageDataGenerator(rescale=1./255)  
test_datagen = ImageDataGenerator(rescale=1./255)
```

# 1. Load and Process the Dataset

```
from keras.preprocessing.image import ImageDataGenerator

# All images will be rescaled by 1./255
train_datagen = ImageDataGenerator(rescale=1./255)
test_datagen = ImageDataGenerator(rescale=1./255)

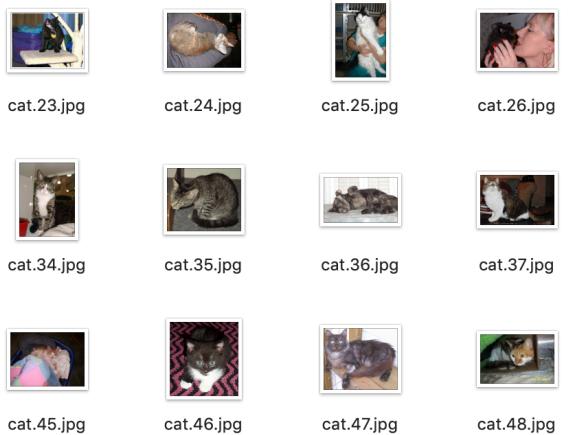
train_generator = train_datagen.flow_from_directory(
    # This is the target directory
    train_dir,
    # All images will be resized to 150x150
    target_size=(150, 150),
    batch_size=20,
    # Since we use binary_crossentropy loss, we need binary labels
    class_mode='binary')
```

# 1. Load and Process the Dataset

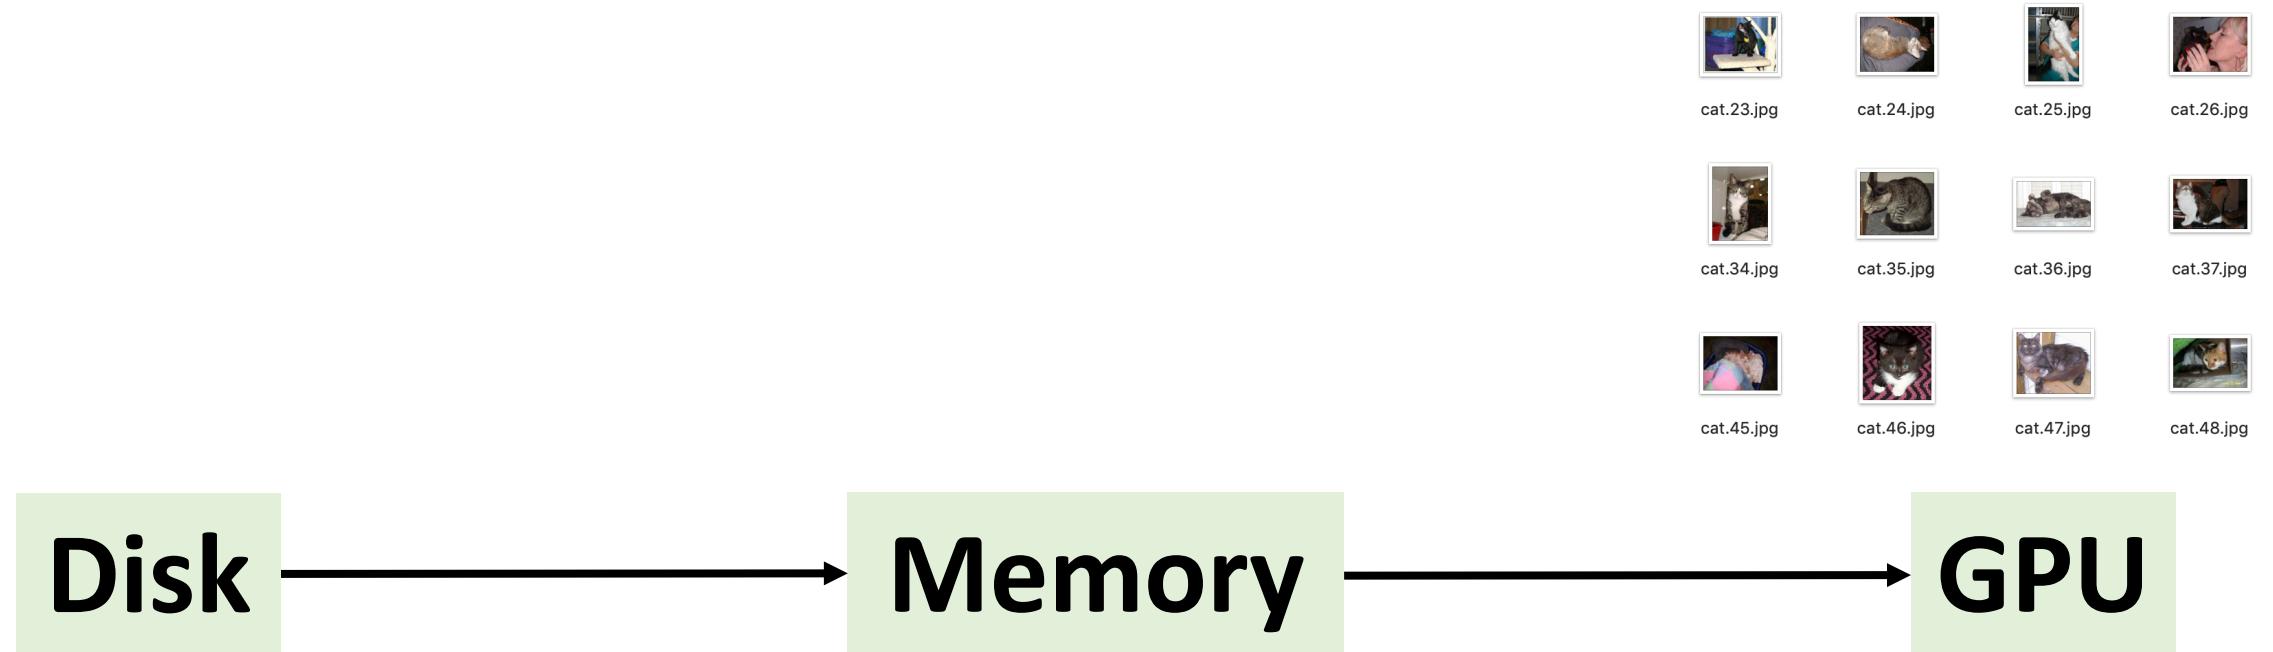
```
for data_batch, labels_batch in train_generator:  
    print('data batch shape:', data_batch.shape)  
    print('labels batch shape:', labels_batch.shape)  
    break
```

```
data batch shape: (20, 150, 150, 3)  
labels batch shape: (20,)
```

# 1. Load and Process the Dataset



# 1. Load and Process the Dataset



## 2. Build the CNN

```
from keras import layers
from keras import models

model = models.Sequential()
model.add(layers.Conv2D(32, (3, 3), activation='relu',
                      input_shape=(150, 150, 3)))
model.add(layers.MaxPooling2D((2, 2)))
model.add(layers.Conv2D(64, (3, 3), activation='relu'))
model.add(layers.MaxPooling2D((2, 2)))
model.add(layers.Conv2D(128, (3, 3), activation='relu'))
model.add(layers.MaxPooling2D((2, 2)))
model.add(layers.Conv2D(128, (3, 3), activation='relu'))
model.add(layers.MaxPooling2D((2, 2)))
model.add(layers.Flatten())
model.add(layers.Dense(512, activation='relu'))
model.add(layers.Dense(1, activation='sigmoid'))
```

## 2. Build the CNN

```
model.summary()
```

Layer (type)	Output Shape	Param #
<hr/>		
conv2d_1 (Conv2D)	(None, 148, 148, 32)	896
max_pooling2d_1 (MaxPooling2D)	(None, 74, 74, 32)	0
conv2d_2 (Conv2D)	(None, 72, 72, 64)	18496
max_pooling2d_2 (MaxPooling2D)	(None, 36, 36, 64)	0
conv2d_3 (Conv2D)	(None, 34, 34, 128)	73856
max_pooling2d_3 (MaxPooling2D)	(None, 17, 17, 128)	0
conv2d_4 (Conv2D)	(None, 15, 15, 128)	147584
max_pooling2d_4 (MaxPooling2D)	(None, 7, 7, 128)	0
flatten_1 (Flatten)	(None, 6272)	0
dense_1 (Dense)	(None, 512)	3211776
dense_2 (Dense)	(None, 1)	513
<hr/>		
Total params: 3,453,121		
Trainable params: 3,453,121		
Non-trainable params: 0		

### 3. Train the CNN

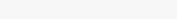
Specify: optimization method, learning rate (LR), loss function, and metric.

```
from keras import optimizers  
  
model.compile(loss='binary_crossentropy',  
                optimizer=optimizers.RMSprop(lr=1e-4),  
                metrics=[ 'acc' ])
```

### 3. Train the CNN

```
history = model.fit_generator(  
    train_generator,  
    steps_per_epoch=100,  
    epochs=30,  
    validation_data=validation_generator,  
    validation_steps=50 )
```

## 3. Train the CNN

```
history = model.fit_generator(  
    train_generator,  
    steps_per_epoch=100,   
    epochs=30,  
    validation_data=validation_generator,  
    validation_steps=50)
```

- Totally  $n = 200$
- Batch size is  $b =$
- Thus  $\frac{n}{b} = 100$  ba

- Totally  $n = 2000$  training samples.
  - Batch size is  $b = 20$ .
  - Thus  $\frac{n}{b} = 100$  batches per epoch.

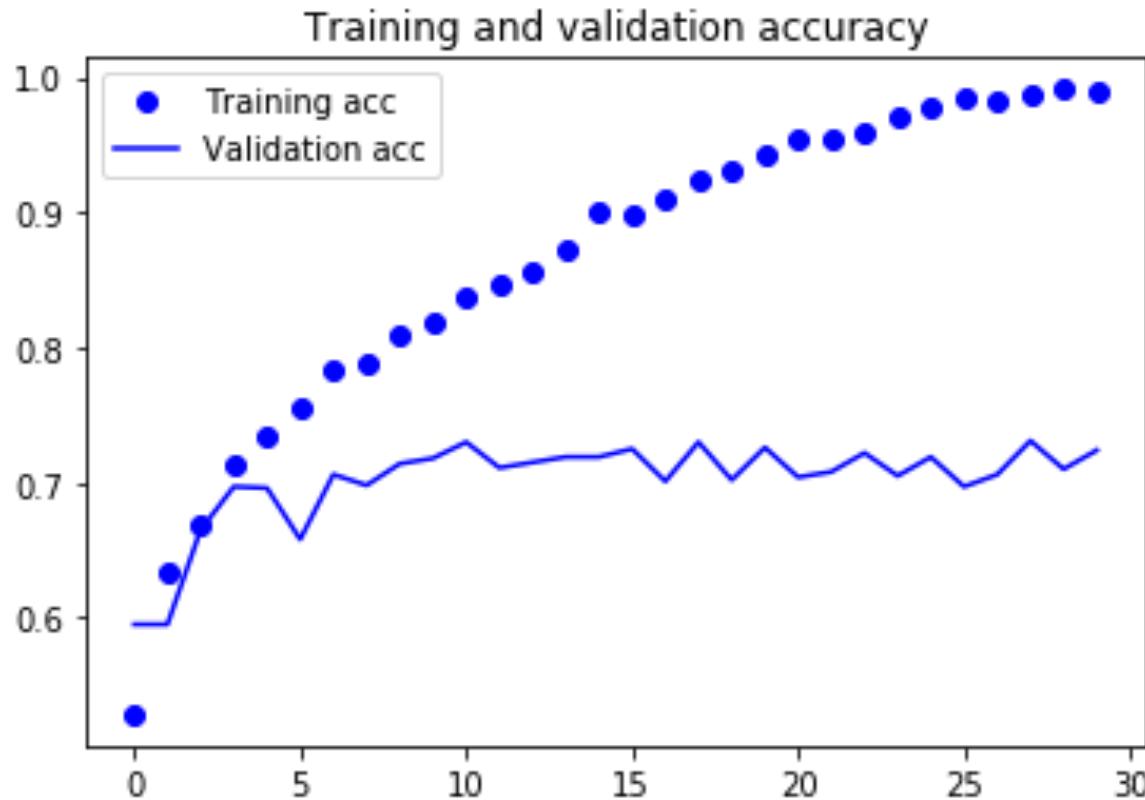
```
Epoch 1/30  
100/100 [=====] - 9s - loss: 0.6898 - acc: 0.5285 - val_loss: 0.6724 - val_acc: 0.5950  
Epoch 2/30  
100/100 [=====] - 8s - loss: 0.6543 - acc: 0.6340 - val_loss: 0.6565 - val_acc: 0.5950  
Epoch 3/30  
100/100 [=====] - 8s - loss: 0.6143 - acc: 0.6690 - val_loss: 0.6116 - val_acc: 0.6650  
Epoch 4/30  
100/100 [=====] - 8s - loss: 0.5626 - acc: 0.7125 - val_loss: 0.5774 - val_acc: 0.6970
```

100

```
Epoch 29/30  
100/100 [=====] - 8s - loss: 0.0375 - acc: 0.9915 - val_loss: 0.9987 - val_acc: 0.7100  
Epoch 30/30  
100/100 [=====] - 8s - loss: 0.0387 - acc: 0.9895 - val_loss: 1.0139 - val_acc: 0.7240
```

# 4. Examine the Results

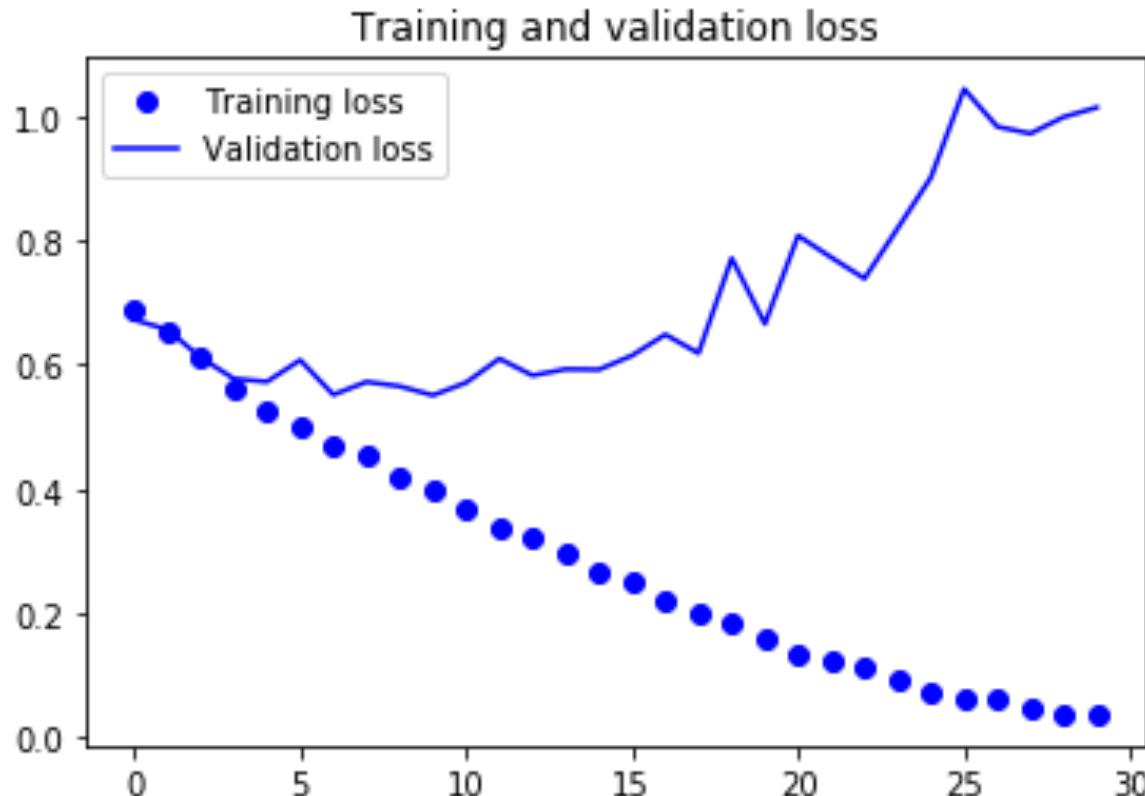
Plot the *accuracy* against *epochs* (1 epoch = 1 pass over the data).



- Training acc: 99.0%
- Validation acc: 72.4%
- Likely overfitting

# 4. Examine the Results

Plot the *loss* against *epochs* (1 epoch = 1 pass over the data).



- Training loss is decreasing.
- Validation loss decreases and then increase.

# Why Overfitting?

```
=====
```

Total params: 3,453,121

Trainable params: 3,453,121

Non-trainable params: 0

---

- Over  $3M$  parameters; Just  $2K$  training samples.
- Overfitting is not surprising.

# Trick 1: Dropout

# Keras's Dropout Layer

- Dropout only before the 1<sup>st</sup> dense layer to regularize the 1<sup>st</sup> dense layer.
- Because the 1<sup>st</sup> dense layer has too many trainable parameters.

```
model = models.Sequential()
model.add(layers.Conv2D(32, (3, 3), activation='relu',
                      input_shape=(150, 150, 3)))
model.add(layers.MaxPooling2D((2, 2)))
model.add(layers.Conv2D(64, (3, 3), activation='relu'))
model.add(layers.MaxPooling2D((2, 2)))
model.add(layers.Conv2D(128, (3, 3), activation='relu'))
model.add(layers.MaxPooling2D((2, 2)))
model.add(layers.Conv2D(128, (3, 3), activation='relu'))
model.add(layers.MaxPooling2D((2, 2)))
model.add(layers.Flatten())
model.add(layers.Dropout(0.5))
model.add(layers.Dense(512, activation='relu'))
model.add(layers.Dense(1, activation='sigmoid'))

model.compile(loss='binary_crossentropy',
              optimizer=optimizers.RMSprop(lr=1e-4),
              metrics=['acc'])
```

# Trick 2: Data Augmentation

# Setup Data Augmentation Using Keras

```
train_datagen = ImageDataGenerator(  
    rescale=1./255,  
    rotation_range=40,  
    width_shift_range=0.2,  
    height_shift_range=0.2,  
    shear_range=0.2,  
    zoom_range=0.2,  
    horizontal_flip=True, )
```

```
train_generator = train_datagen.flow_from_directory(  
    # This is the target directory  
    train_dir,  
    # All images will be resized to 150x150  
    target_size=(150, 150),  
    batch_size=32,  
    # Since we use binary_crossentropy loss, we need binary labels  
    class_mode='binary')
```

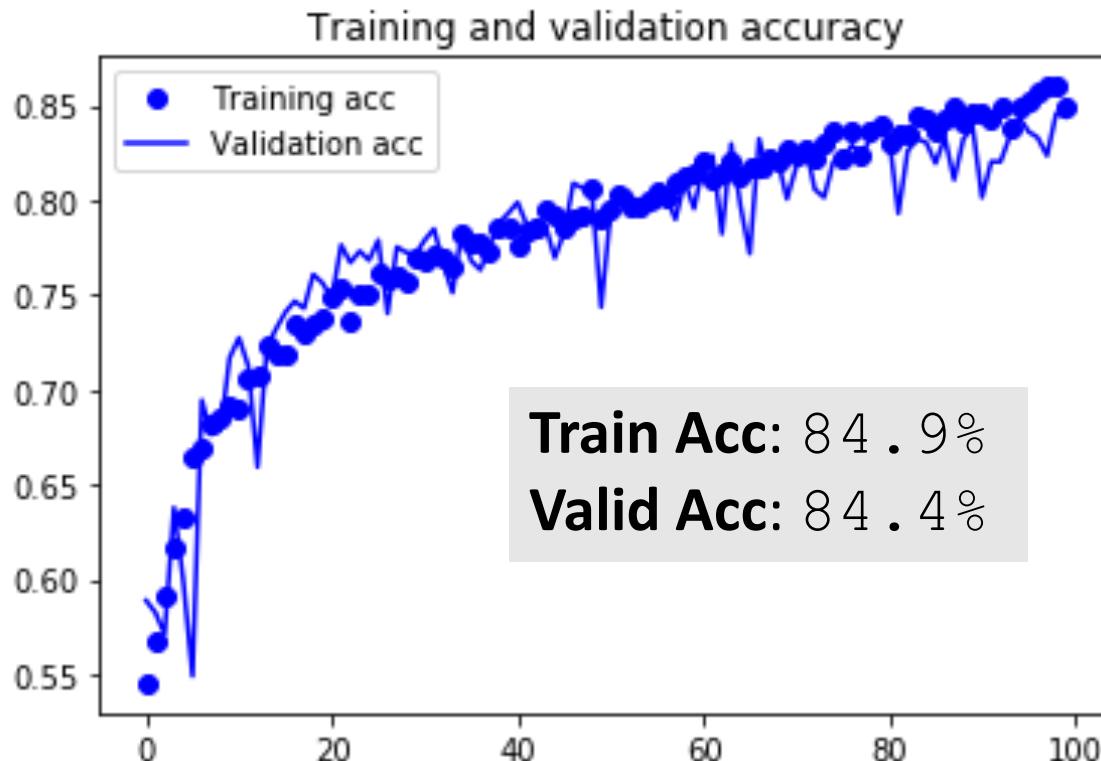
# Train the CNN

```
history = model.fit_generator(  
    train_generator,  
    steps_per_epoch=100,  
    epochs=100,  
    validation_data=validation_generator,  
    validation_steps=50)
```

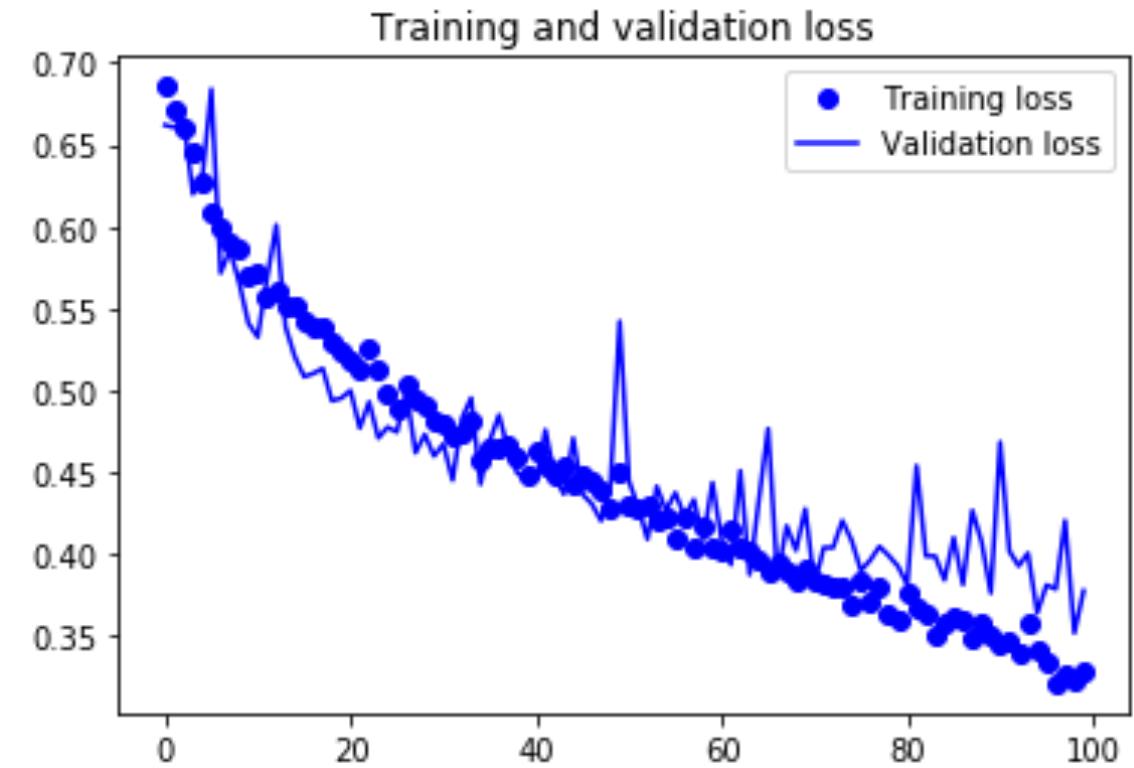
```
Epoch 1/100  
100/100 [=====] - 24s - loss: 0.6857 - acc: 0.5447 - val_loss: 0.6620 - val_acc: 0.5888  
Epoch 2/100  
100/100 [=====] - 23s - loss: 0.6710 - acc: 0.5675 - val_loss: 0.6606 - val_acc: 0.5825  
Epoch 3/100  
100/100 [=====] - 22s - loss: 0.6609 - acc: 0.5913 - val_loss: 0.6663 - val_acc: 0.5711  
●  
●  
●  
Epoch 99/100  
100/100 [=====] - 22s - loss: 0.3255 - acc: 0.8581 - val_loss: 0.3518 - val_acc: 0.8460  
Epoch 100/100  
100/100 [=====] - 22s - loss: 0.3280 - acc: 0.8491 - val_loss: 0.3776 - val_acc: 0.8439
```

# Examine the Results

*accuracy against epochs*



*loss against epochs*



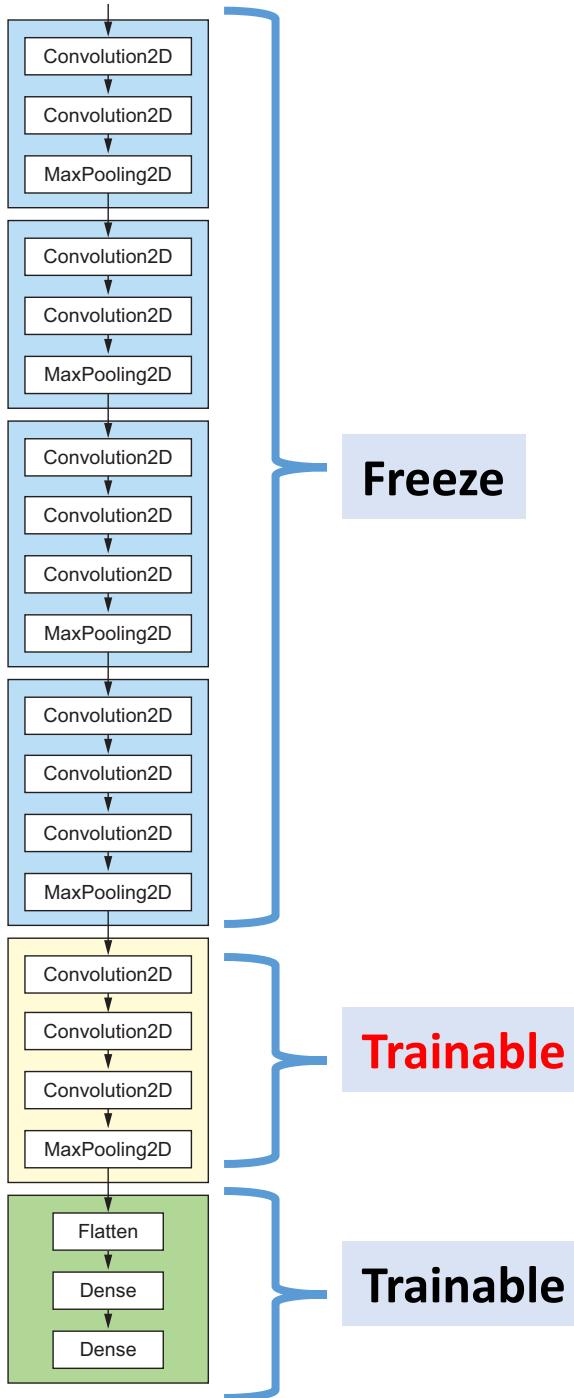
# Trick 3: Pretrain

# Train a Deep Neural Network?

```
model = models.Sequential()
model.add(layers.Conv2D(32, (3, 3), activation='relu',
                      input_shape=(150, 150, 3)))
model.add(layers.MaxPooling2D((2, 2)))
model.add(layers.Conv2D(64, (3, 3), activation='relu'))
model.add(layers.MaxPooling2D((2, 2)))
model.add(layers.Conv2D(128, (3, 3), activation='relu'))
model.add(layers.MaxPooling2D((2, 2)))
model.add(layers.Conv2D(128, (3, 3), activation='relu'))
model.add(layers.MaxPooling2D((2, 2)))
model.add(layers.Flatten())
model.add(layers.Dropout(0.5))
model.add(layers.Dense(512, activation='relu'))
model.add(layers.Dense(1, activation='sigmoid'))

model.compile(loss='binary_crossentropy',
              optimizer=optimizers.RMSprop(lr=1e-4),
              metrics=[ 'acc' ])
```

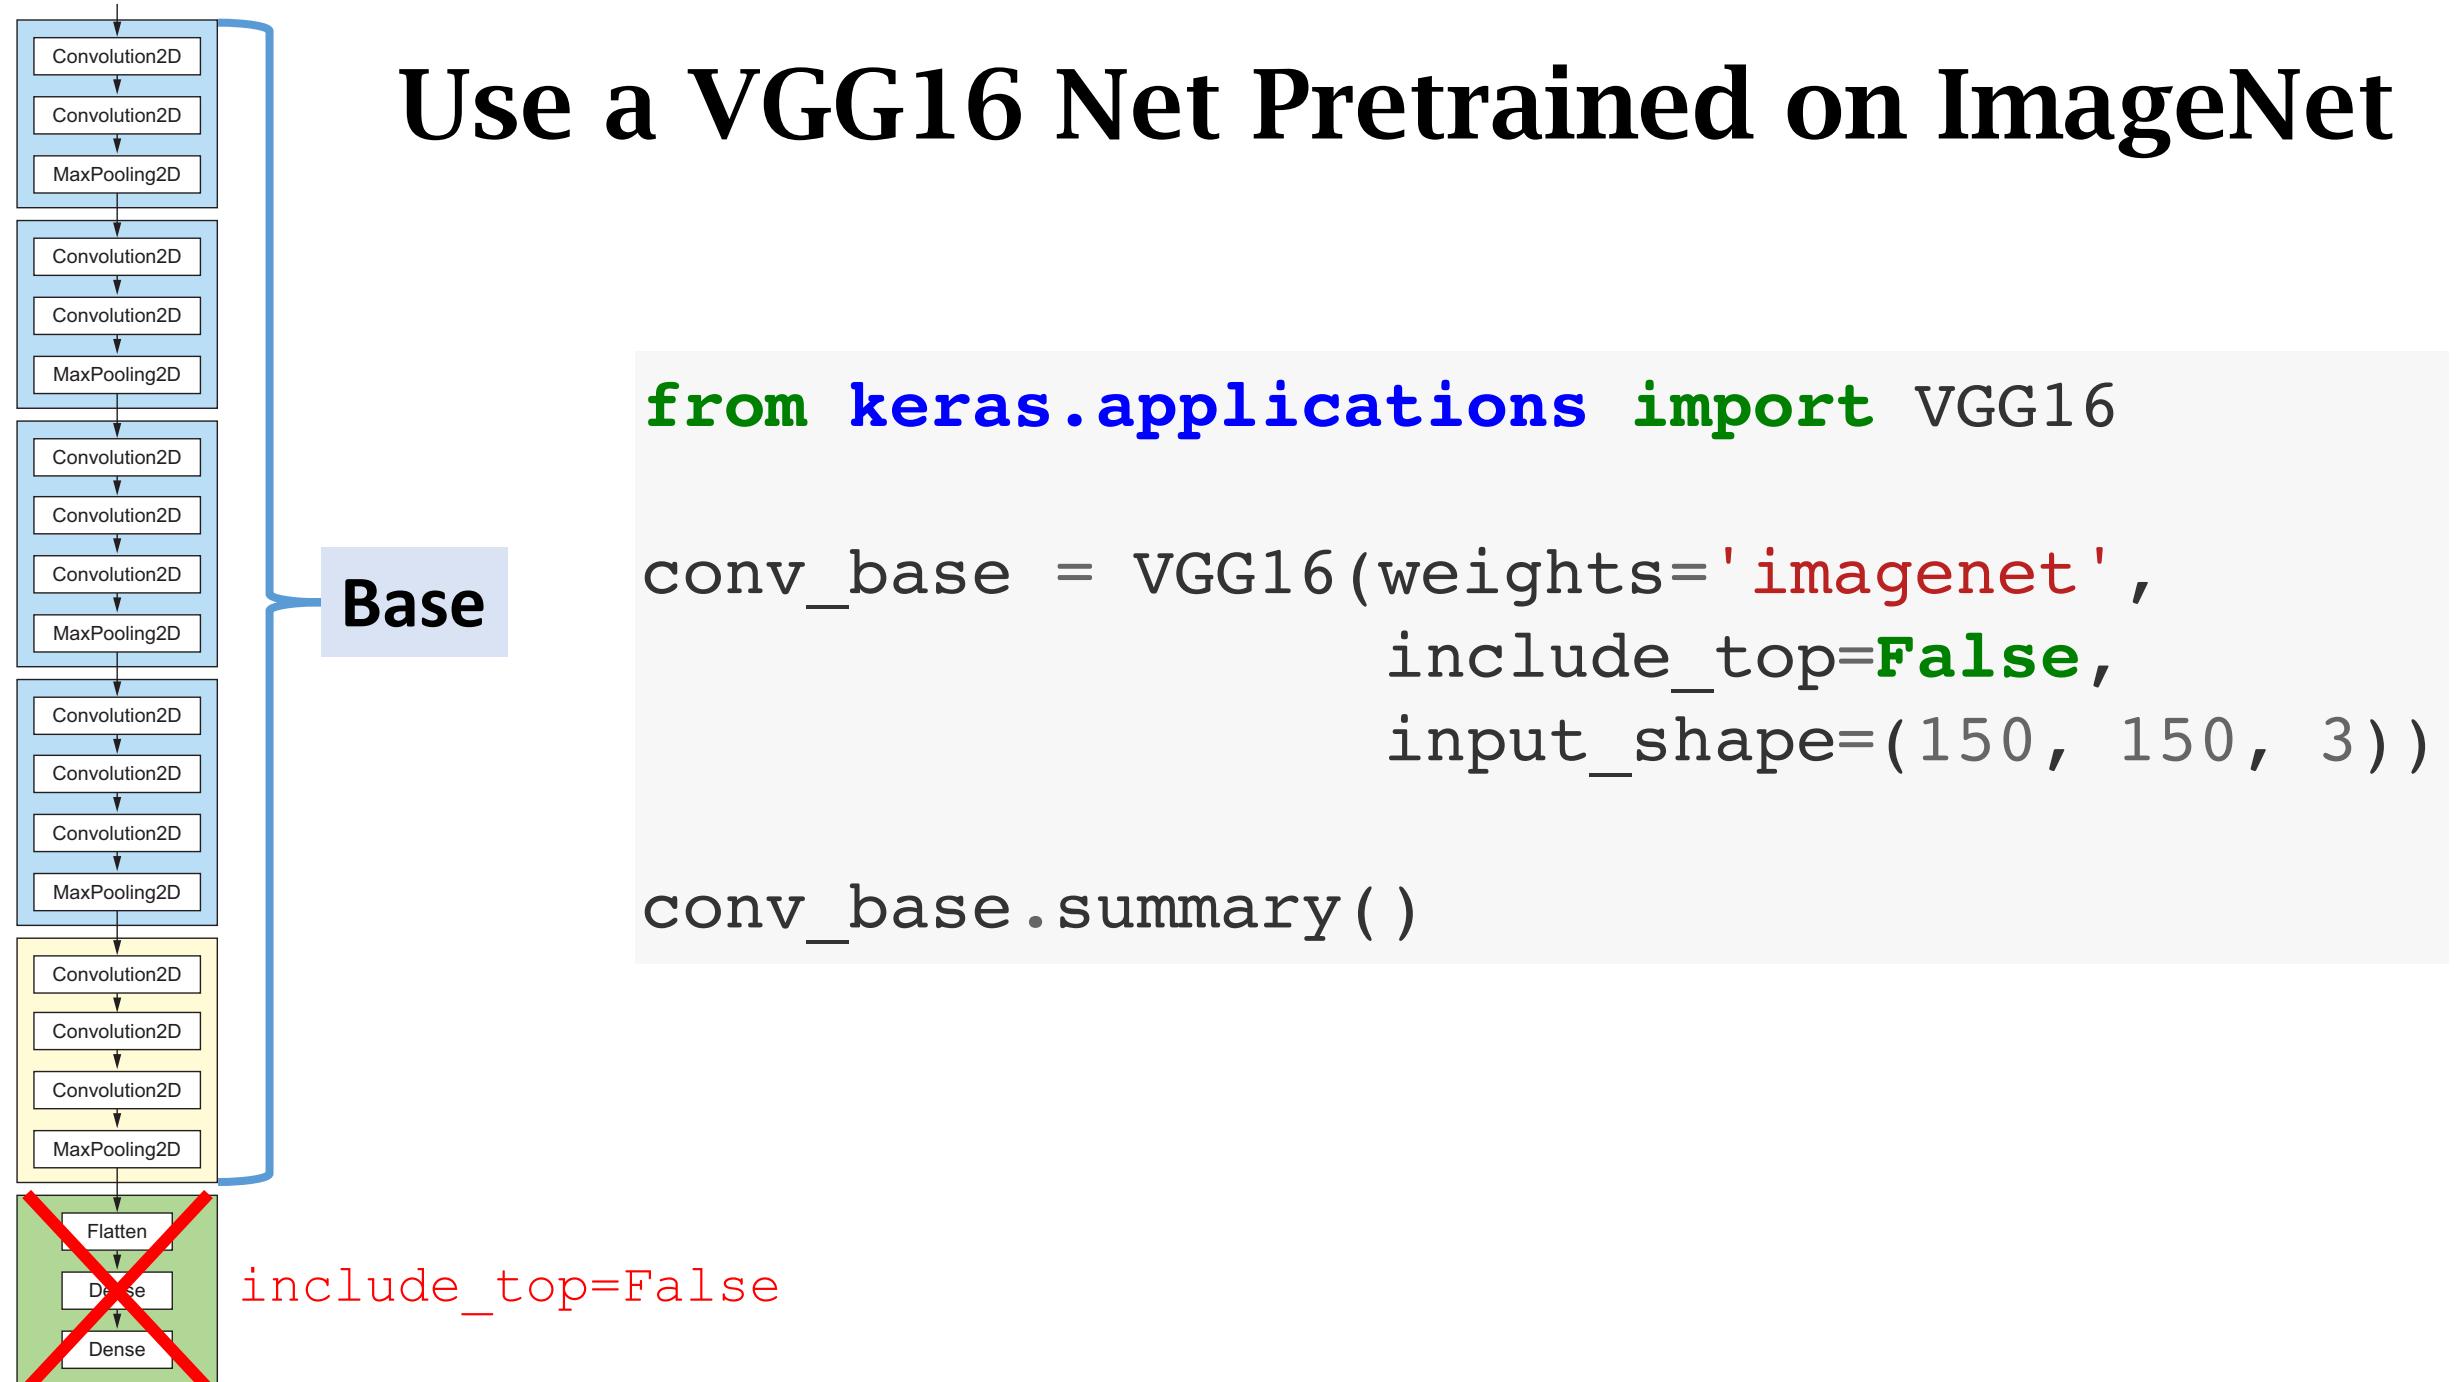
- We have trained a neural net with 4 Conv Layers and 2 FC Layers.
- Relatively shallow.

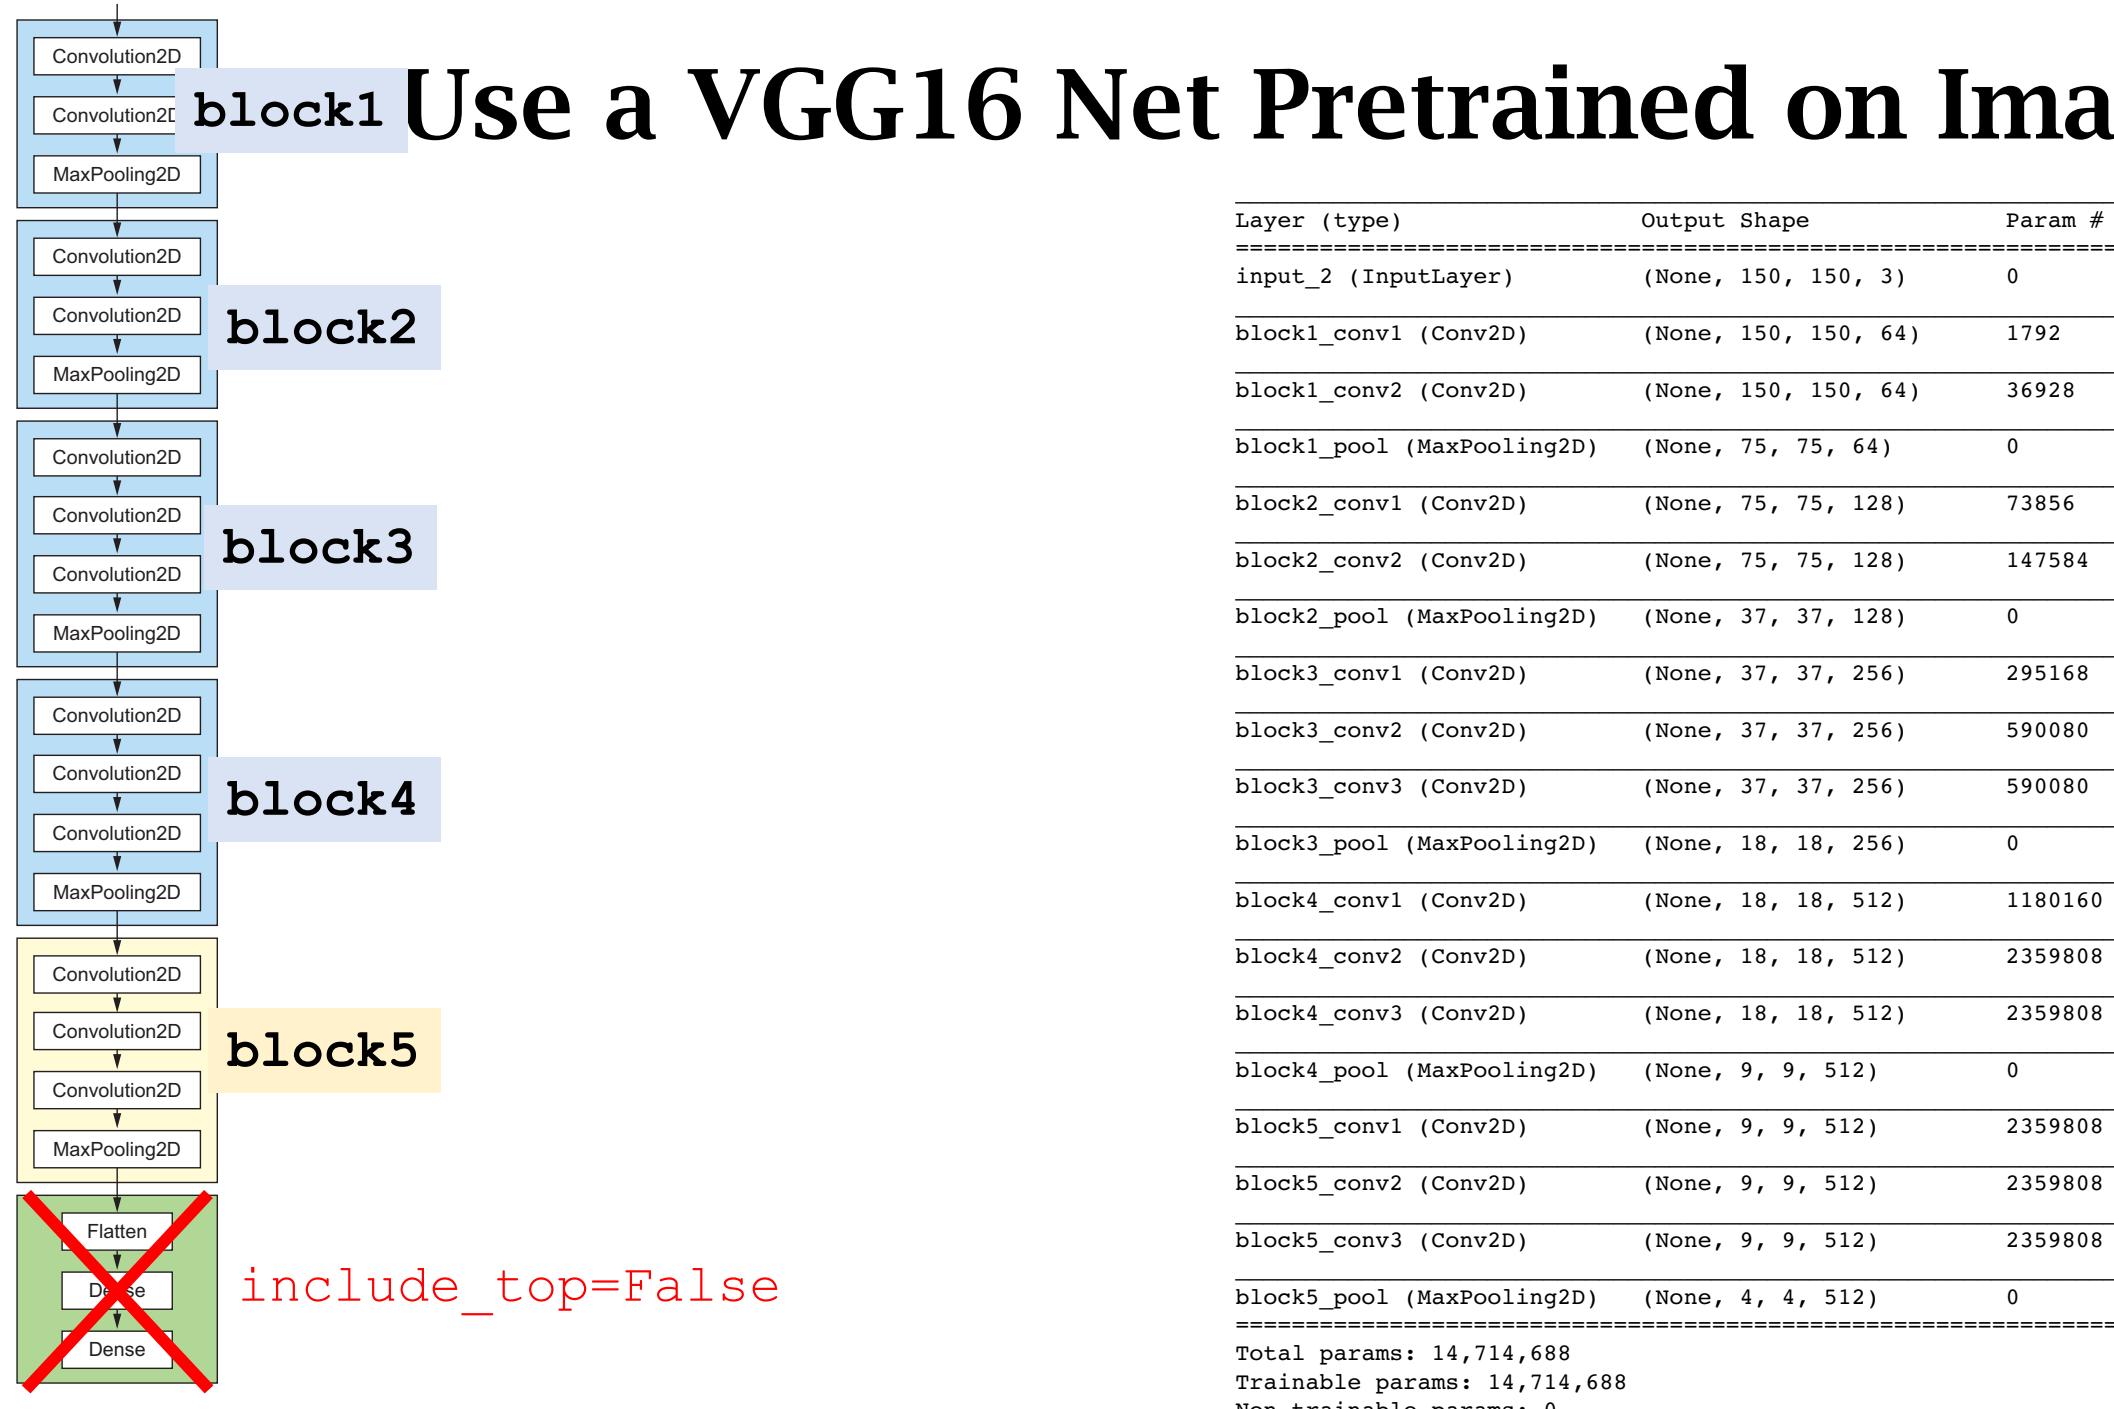


# Pretrain

1. Pre-train a deep net on large-scale dataset, e.g., ImageNet ( $14M$  images with labels).
2. Remove the top layers.
3. Build new top layers (randomly initialized).
4. Freeze the base layers; Train the top layers.
5. Optional: **Fine-tune** the top Conv Layers.

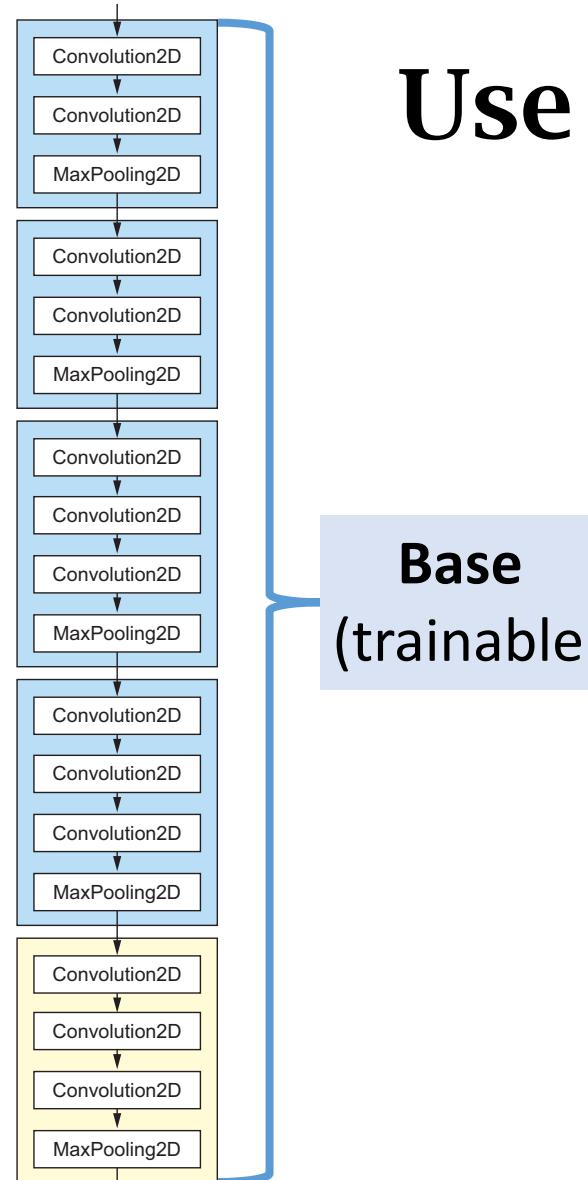
# Use a VGG16 Net Pretrained on ImageNet





Layer (type)	Output Shape	Param #
input_2 (InputLayer)	(None, 150, 150, 3)	0
block1_conv1 (Conv2D)	(None, 150, 150, 64)	1792
block1_conv2 (Conv2D)	(None, 150, 150, 64)	36928
block1_pool (MaxPooling2D)	(None, 75, 75, 64)	0
block2_conv1 (Conv2D)	(None, 75, 75, 128)	73856
block2_conv2 (Conv2D)	(None, 75, 75, 128)	147584
block2_pool (MaxPooling2D)	(None, 37, 37, 128)	0
block3_conv1 (Conv2D)	(None, 37, 37, 256)	295168
block3_conv2 (Conv2D)	(None, 37, 37, 256)	590080
block3_conv3 (Conv2D)	(None, 37, 37, 256)	590080
block3_pool (MaxPooling2D)	(None, 18, 18, 256)	0
block4_conv1 (Conv2D)	(None, 18, 18, 512)	1180160
block4_conv2 (Conv2D)	(None, 18, 18, 512)	2359808
block4_conv3 (Conv2D)	(None, 18, 18, 512)	2359808
block4_pool (MaxPooling2D)	(None, 9, 9, 512)	0
block5_conv1 (Conv2D)	(None, 9, 9, 512)	2359808
block5_conv2 (Conv2D)	(None, 9, 9, 512)	2359808
block5_conv3 (Conv2D)	(None, 9, 9, 512)	2359808
block5_pool (MaxPooling2D)	(None, 4, 4, 512)	0
<hr/>		
Total params: 14,714,688		
Trainable params: 14,714,688		
Non-trainable params: 0		

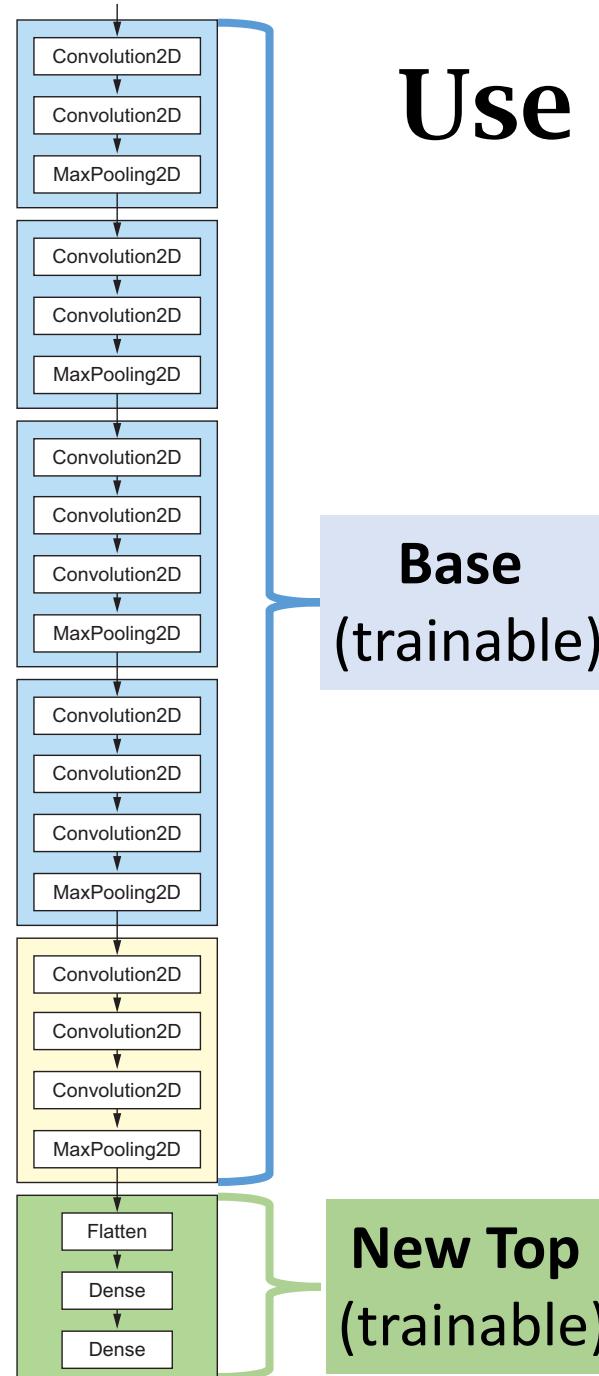
# Use a VGG16 Net Pretrained on ImageNet



```
from keras import models
from keras import layers

model = models.Sequential()
model.add(conv_base)
```

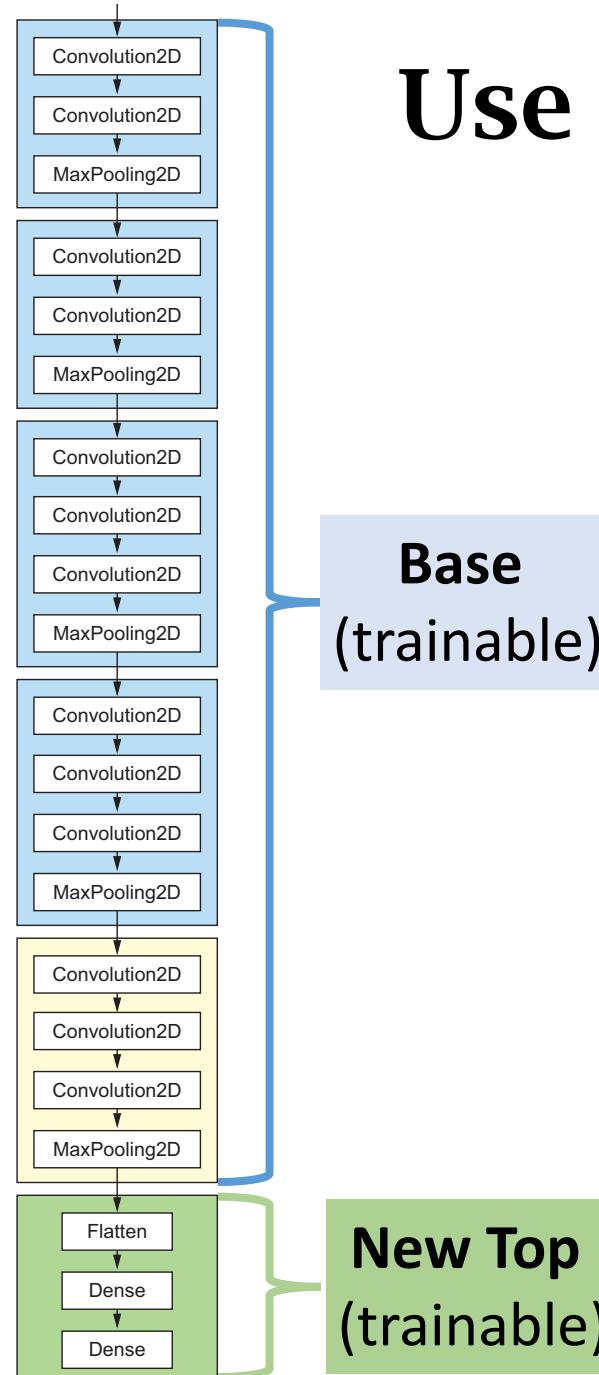
# Use a VGG16 Net Pretrained on ImageNet



```
from keras import models
from keras import layers

model = models.Sequential()
model.add(conv_base)
model.add(layers.Flatten())
model.add(layers.Dense(256, activation='relu'))
model.add(layers.Dense(1, activation='sigmoid'))
```

# Use a VGG16 Net Pretrained on ImageNet



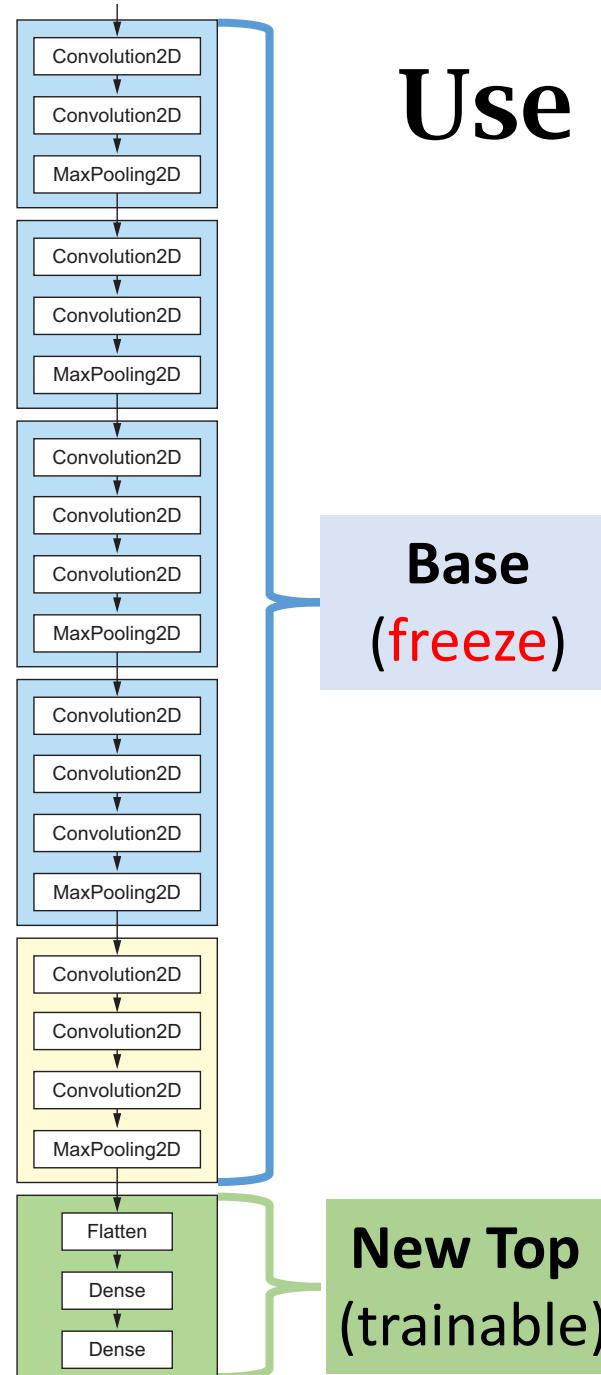
Layer (type)	Output Shape	Param #
vgg16 (Model)	(None, 4, 4, 512)	14714688
flatten_1 (Flatten)	(None, 8192)	0
dense_1 (Dense)	(None, 256)	2097408
dense_2 (Dense)	(None, 1)	257

Total params: 16,812,353

Trainable params: 16,812,353

Non-trainable params: 0

# Use a VGG16 Net Pretrained on ImageNet



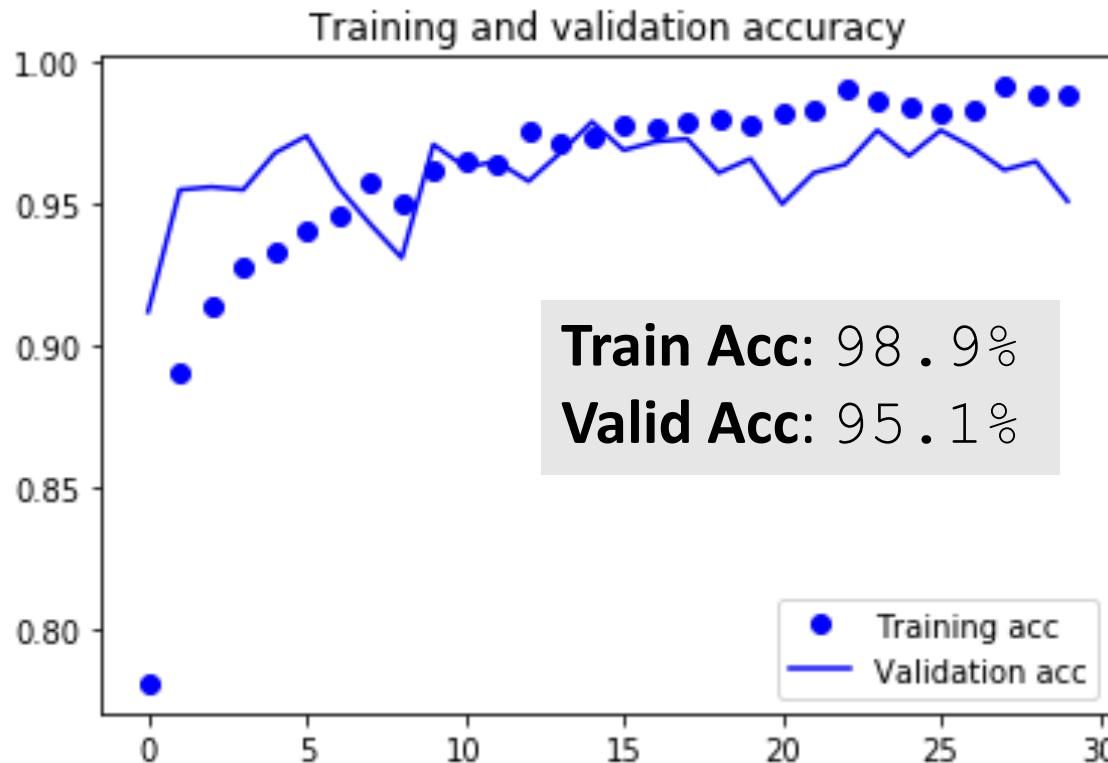
```
conv_base.trainable = False
```

```
model.summary()
```

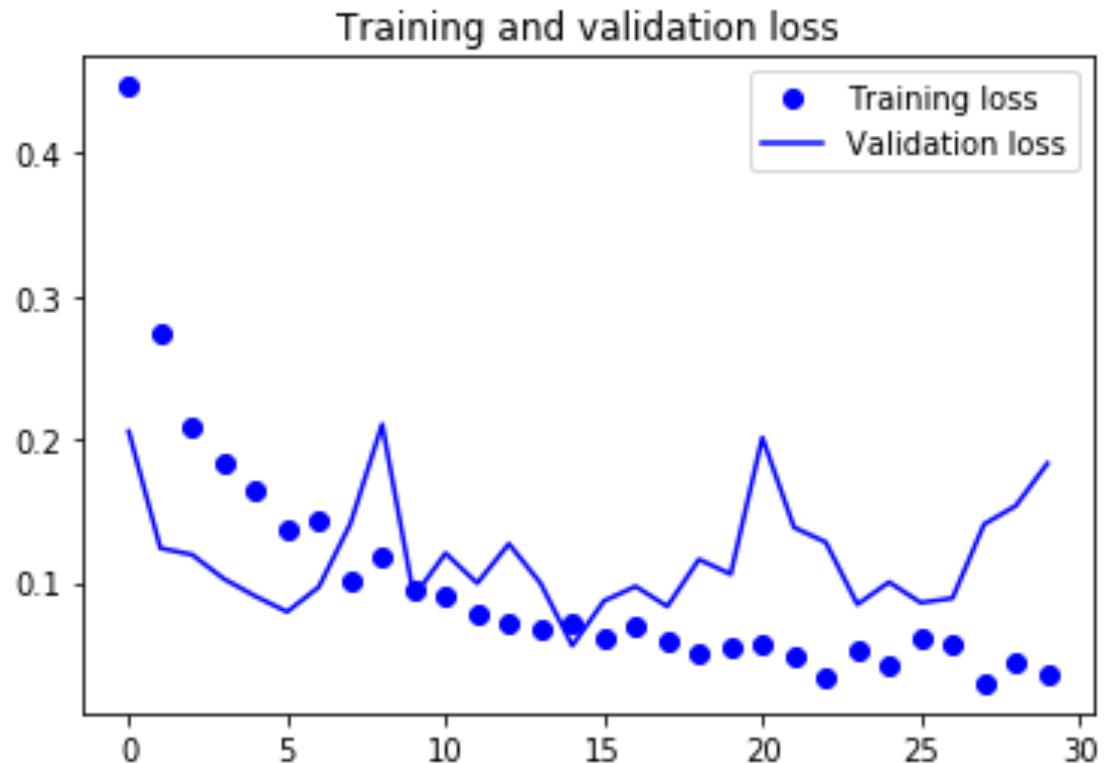
Layer (type)	Output Shape	Param #
vgg16 (Model)	(None, 4, 4, 512)	14714688
flatten_1 (Flatten)	(None, 8192)	0
dense_1 (Dense)	(None, 256)	2097408
dense_2 (Dense)	(None, 1)	257
Total params: 16,812,353		
Trainable params: 2,097,665		
Non-trainable params: 14,714,688		

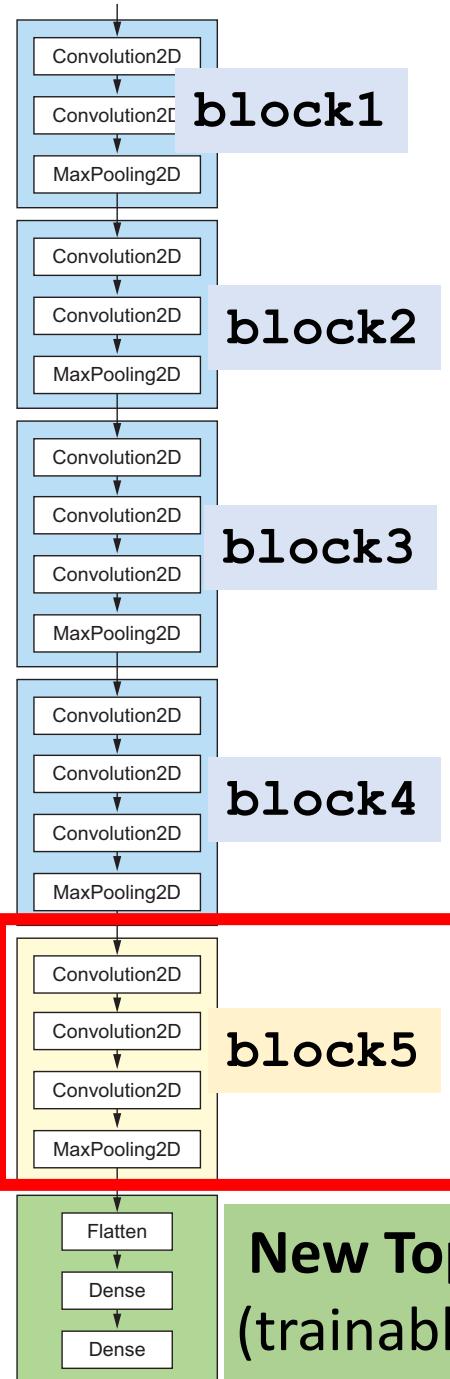
# After Training the New Top

*accuracy* against *epochs*



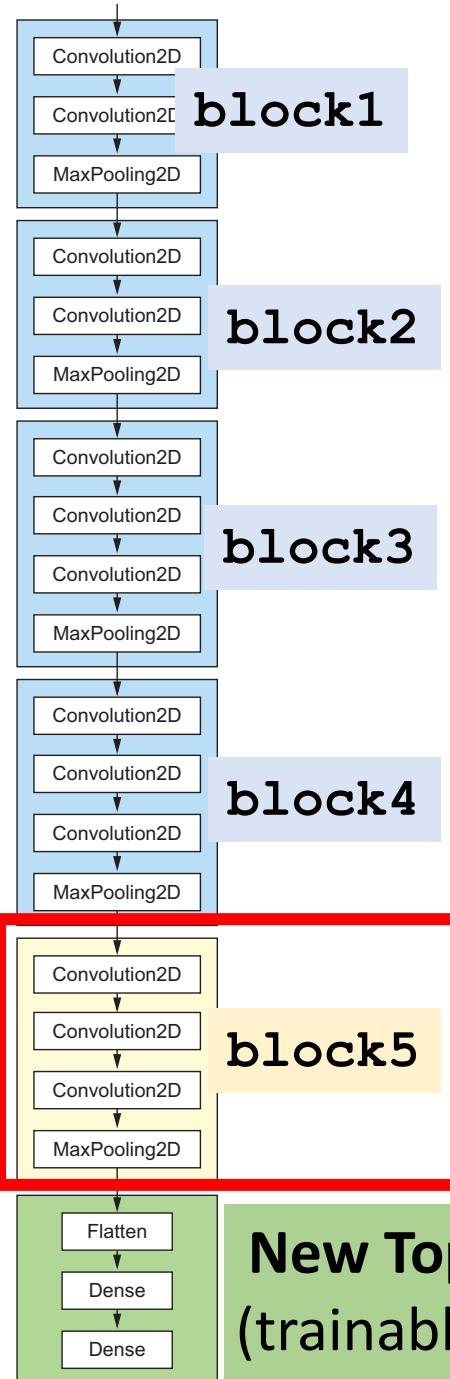
*loss* against *epochs*





# Fine Tuning the Top Conv Layers

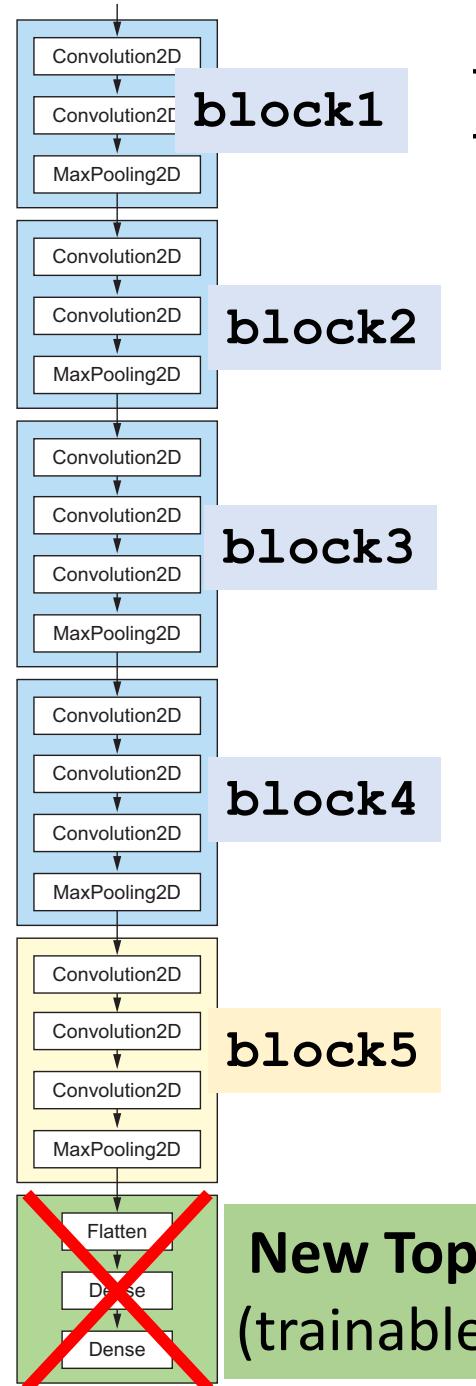
Layer (type)	Output Shape	Param #
input_2 (InputLayer)	(None, 150, 150, 3)	0
block1_conv1 (Conv2D)	(None, 150, 150, 64)	1792
block1_conv2 (Conv2D)	(None, 150, 150, 64)	36928
block1_pool (MaxPooling2D)	(None, 75, 75, 64)	0
block2_conv1 (Conv2D)	(None, 75, 75, 128)	73856
block2_conv2 (Conv2D)	(None, 75, 75, 128)	147584
block2_pool (MaxPooling2D)	(None, 37, 37, 128)	0
block3_conv1 (Conv2D)	(None, 37, 37, 256)	295168
block3_conv2 (Conv2D)	(None, 37, 37, 256)	590080
block3_conv3 (Conv2D)	(None, 37, 37, 256)	590080
block3_pool (MaxPooling2D)	(None, 18, 18, 256)	0
block4_conv1 (Conv2D)	(None, 18, 18, 512)	1180160
block4_conv2 (Conv2D)	(None, 18, 18, 512)	2359808
block4_conv3 (Conv2D)	(None, 18, 18, 512)	2359808
block4_pool (MaxPooling2D)	(None, 9, 9, 512)	0
block5_conv1 (Conv2D)	(None, 9, 9, 512)	2359808
block5_conv2 (Conv2D)	(None, 9, 9, 512)	2359808
block5_conv3 (Conv2D)	(None, 9, 9, 512)	2359808
block5_pool (MaxPooling2D)	(None, 4, 4, 512)	0
<hr/>		
Total params: 14,714,688		
Trainable params: 14,714,688		
Non-trainable params: 0		



# Fine Tuning the Top Conv Layers

```
trainable_layer_names = ['block5_conv1', 'block5_conv2',
                         'block5_conv3', 'block5_pool']
conv_base.trainable = True

for layer in conv_base.layers:
    if layer.name in trainable_layer_names:
        layer.trainable = True
    else:
        layer.trainable = False
```



# Fine Tuning the Top Conv Layers

```
model.summary()
```

Layer (type)	Output Shape	Param #
vgg16 (Model)	(None, 4, 4, 512)	14714688
flatten_1 (Flatten)	(None, 8192)	0
dense_1 (Dense)	(None, 256)	2097408
dense_2 (Dense)	(None, 1)	257
<hr/>		
Total params: 16,812,353		
Trainable params: 9,177,089		
Non-trainable params: 7,635,264		

# Fine Tuning the Top Conv Layers

**Re-compile before training**

```
model.compile(loss='binary_crossentropy',  
              optimizer=optimizers.RMSprop(lr=1e-5),  
              metrics=[ 'acc' ] )
```

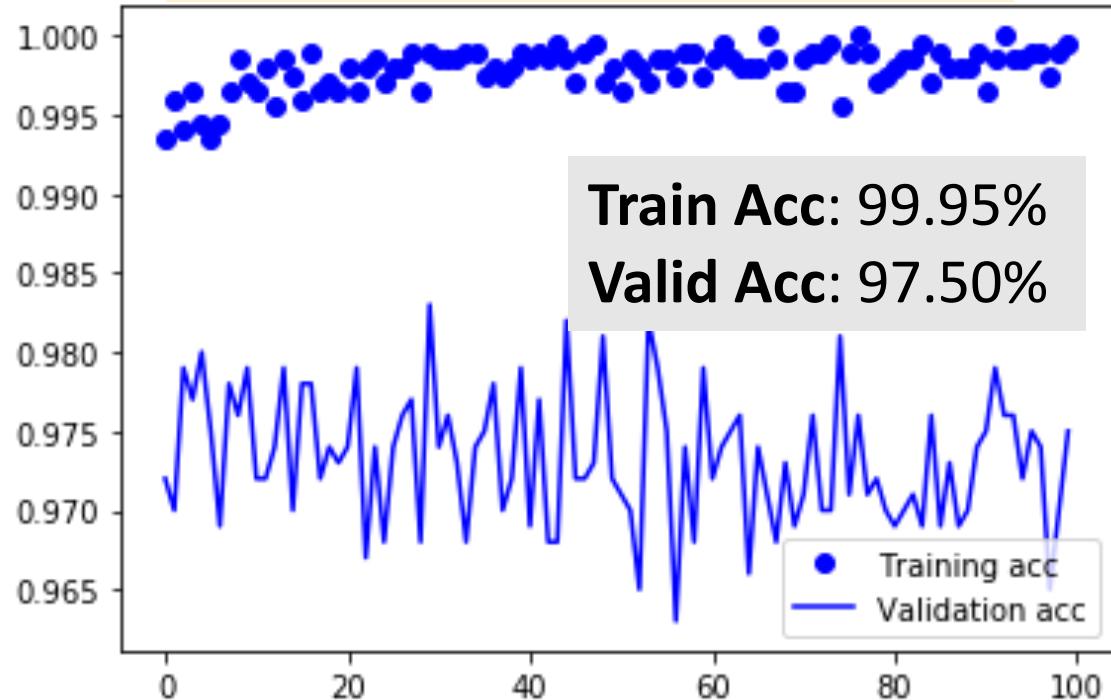
# Fine Tuning the Top Conv Layers

```
history = model.fit_generator(  
    train_generator,  
    steps_per_epoch=100,  
    epochs=100,  
    validation_data=validation_generator,  
    validation_steps=50)
```

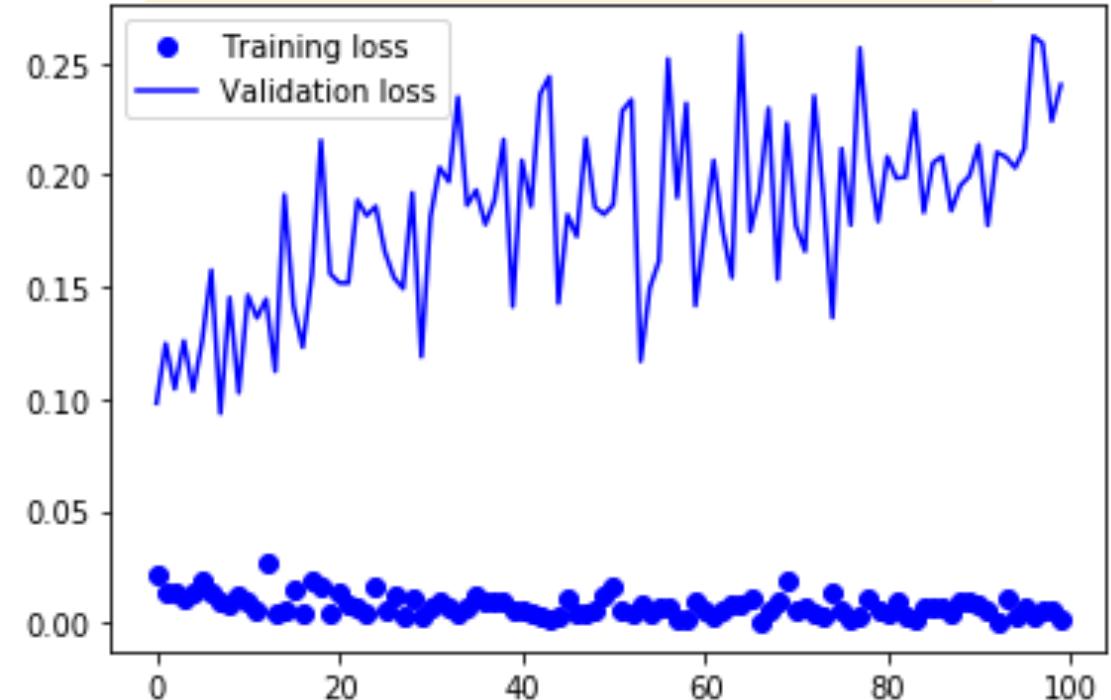
```
Epoch 1/100  
100/100 [=====] - 32s - loss: 0.0215 - acc: 0.9935 - val_loss: 0.0980 - val_acc: 0.9720  
Epoch 2/100  
100/100 [=====] - 32s - loss: 0.0131 - acc: 0.9960 - val_loss: 0.1247 - val_acc: 0.9700  
Epoch 3/100  
100/100 [=====] - 32s - loss: 0.0140 - acc: 0.9940 - val_loss: 0.1044 - val_acc: 0.9790  
•  
•  
•  
Epoch 99/100  
100/100 [=====] - 33s - loss: 0.0060 - acc: 0.9990 - val_loss: 0.2242 - val_acc: 0.9700  
Epoch 100/100  
100/100 [=====] - 33s - loss: 0.0010 - acc: 0.9995 - val_loss: 0.2403 - val_acc: 0.9750
```

# Fine Tuning the Top Conv Layers

*accuracy* against *epochs*



*loss* against *epochs*



# Fine Tuning the Top Conv Layers

Evaluate the model on the test set

```
test_generator = test_datagen.flow_from_directory(  
    test_dir,  
    target_size=(150, 150),  
    batch_size=20,  
    class_mode='binary')  
  
test_loss, test_acc = model.evaluate_generator(test_generator, steps=50)  
print('test acc:', test_acc)
```

Found 1000 images belonging to 2 classes.  
test acc: 0.967999992371

# Summary of the Results

- A small ConvNet (4 Conv Layers + 2 FC Layers), with  $3.5M$  parameters.
  - Training accuracy: 99.0%
  - Validation accuracy: 72.4%
- The small ConvNet + 1 Dropout Layer + Data Augmentation.
  - Training accuracy: 84.9%
  - Validation accuracy: 84.4%
- VGG16 Net pre-trained on the ImageNet. (Train the new top layers.)
  - Training accuracy: 98.9%
  - Validation accuracy: 95.1%
- VGG16 Net pre-trained on the ImageNet. (Fine-tune the top Conv Layers.)
  - Training accuracy: 99.95%
  - Validation accuracy: 97.5%